


# node2vec: Scalable Feature Learning for Networks

*Aditya Grover and Jure Leskovec. KDD 2016.*

Presented by Haoxiang Wang. Feb 26, 2020.


# Node Embeddings



► **Intuition:** Find embeddings of nodes in a  $d$ -dimensional space so that “similar” nodes in the graph have embeddings that are close together.

# Setup

- ▶ Assume we have a graph  $G$ :
  - ▶  $V$  is the vertex set (i.e., node set).
  - ▶  $A$  is the adjacency matrix (assume binary).



$$\begin{pmatrix} 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{pmatrix}$$


$$\begin{pmatrix} 0 & 1 & 0 & 1 \\ 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 \\ 1 & 0 & 1 & 0 \end{pmatrix}$$

$$\begin{pmatrix} 0 & 1 & 1 & 1 \\ 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{pmatrix}$$

# Embedding Nodes

- Goal: to encode nodes so that **similarity** in the embedding space (e.g., dot product) approximates **similarity** in the original network.



$$\text{similarity}(u, v) \approx \mathbf{z}_v^\top \mathbf{z}_u$$



# Random Walk Embeddings: Basic Idea


$$\mathbf{z}_u^T \mathbf{z}_v \approx \text{probability that } u \text{ and } v \text{ co-occur on a random walk over the network}$$

1. Estimate probability of visiting node  $v$  on a random walk starting from node  $u$  using some random walk strategy  $R$ .
2. Optimize embeddings to encode these random walk statistics.



# Algorithm/Optimization of Random Walk Embeddings


1. Run short random walks starting from each node on the graph using some strategy  $R$ .
2. For each node  $u$  collect  $N_R(u)$ , the multiset\* of nodes visited on random walks starting from  $u$ . (\*  $N_R(u)$  can have repeat elements since nodes can be visited multiple times on random walks.)
3. Optimize embeddings to according to:



$$\mathcal{L} = \sum_{u \in V} \sum_{v \in N_R(u)} -\log(P(v|\mathbf{z}_u))$$


$$P(v|\mathbf{z}_u) = \frac{\exp(\mathbf{z}_u^\top \mathbf{z}_v)}{\sum_{n \in V} \exp(\mathbf{z}_u^\top \mathbf{z}_n)}$$

In practice, random sampling based on some distribution over nodes



# Node2vec: Biased Random Walks

- **Idea:** use flexible, biased random walks that can trade off between **local** and **global** views of the network ([Grover and Leskovec, 2016](#)).
- BFS (Breath-First Search) and DFS (Depth-First Search): Two classic strategies to define a neighborhood  $N_R(u)$  of a given node  $u$ :



$$N_{BFS}(u) = \{s_1, s_2, s_3\}$$

Local microscopic view

$$N_{DFS}(u) = \{s_4, s_5, s_6\}$$


Global macroscopic view

# Combine BFS + DFS by a Ratio

Biased random walk  $R$  that given a node  $u$  generates neighborhood  $N_R(u)$

- Two parameters:
  - Return parameter  $p$ : Return back to the previous node
  - Walk-away parameter  $q$  : Moving outwards (DFS) vs. inwards (BFS)

Walker is at  $w$ . Where to go next?



$s_1$	$1/p$
$s_2$	1
$s_3$	$1/q$


Unnormalized transition prob.

**BFS-like** walk: Low value of  $p$   
**DFS-like** walk: Low value of  $q$

# Benchmarks: Node Classification & Link Prediction

Node  
Classification

Link Prediction



## Link Prediction

# Empirical Results

### Node Classification

Algorithm	Dataset		
	BlogCatalog	PPI	Wikipedia
Spectral Clustering	0.0405	0.0681	0.0395
DeepWalk	0.2110	0.1768	0.1274
LINE	0.0784	0.1447	0.1164
<i>node2vec</i>	<b>0.2581</b>	<b>0.1791</b>	<b>0.1552</b>
<i>node2vec</i> settings (p,q)	0.25, 0.25	4, 1	4, 0.5
<b>Gain of <i>node2vec</i> [%]</b>	<b>22.3</b>	<b>1.3</b>	<b>21.8</b>


Table 2: Macro-F<sub>1</sub> scores for multilabel classification on BlogCatalog, PPI (*Homo sapiens*) and Wikipedia word cooccurrence networks with 50% of the nodes labeled for training.

Op	Algorithm	Dataset		
		Facebook	PPI	arXiv
(a)	Common Neighbors	0.8100	0.7142	0.8153
	Jaccard's Coefficient	0.8880	0.7018	0.8067
	Adamic-Adar	0.8289	0.7126	0.8315
	Pref. Attachment	0.7137	0.6670	0.6996
(b)	Spectral Clustering	0.5960	0.6588	0.5812
	DeepWalk	0.7238	0.6923	0.7066
	LINE	0.7029	0.6330	0.6516
	<i>node2vec</i>	0.7266	0.7543	0.7221
(c)	Spectral Clustering	0.6192	0.4920	0.5740
	DeepWalk	<b>0.9680</b>	0.7441	0.9340
	LINE	0.9490	0.7249	0.8902
	<i>node2vec</i>	<b>0.9680</b>	<b>0.7719</b>	<b>0.9366</b>
(d)	Spectral Clustering	0.7200	0.6356	0.7099
	DeepWalk	0.9574	0.6026	0.8282
	LINE	0.9483	0.7024	0.8809
	<i>node2vec</i>	0.9602	0.6292	0.8468

Table 4: Area Under Curve (AUC) scores for link prediction. Comparison with popular baselines and embedding based methods bootstrapped using binary operators: (a) Average, (b) Hadamard, (c) Weighted-L1, and (d) Weighted-L2 (See Table 1 for definitions).


# Advantages of Node2Vec

- ▶ node2vec performs better on **node classification** compared with other node embedding methods.
- ▶ Random walk approaches are generally more efficient (i.e.,  $O(|E|)$  vs.  $O(|V|^2)$ )
- ▶ (Note: In general, one must choose definition of node similarity that matches application. )



# Other random walk node embedding works

- ▶ **Different kinds of biased random walks:**
  - ▶ Based on node attributes ([Dong et al., 2017](#)).
  - ▶ Based on a learned weights ([Abu-El-Haija et al., 2017](#))
- ▶ **Alternative optimization schemes:**
  - ▶ Directly optimize based on 1-hop and 2-hop random walk probabilities (as in [LINE from Tang et al. 2015](#)).
- ▶ **Network preprocessing techniques:**
  - ▶ Run random walks on modified versions of the original network (e.g., [Ribeiro et al. 2017's struct2vec](#), [Chen et al. 2016's HARP](#)).




---

# LEARNING ENTITY AND RELATION EMBEDDINGS FOR KNOWLEDGE GRAPH COMPLETION

XIAODAN DU

# KNOWLEDGE GRAPH COMPLETION

- Predicting relations between entities under supervision of the existing knowledge graph




## KNOWLEDGE GRAPH EMBEDDING

- Embedding a knowledge graph into a continuous vector space while preserving certain information of the graph
- Learning vector embeddings for both entities and relationships
  - TransE (Bordes et al. 2013), TransH (Wang et al. 2014): assume embeddings of entities and relations belong to a single space  $\mathbb{R}^k$
  - TransR: assumes one entity space and multiple relation spaces

# TRANSE AND TRANSH


IF TRIPLE (H, R, T) HOLDS

TransE



$$f_r(h, t) = \|\mathbf{h} + \mathbf{r} - \mathbf{t}\|_2^2$$

TransH



$$f_r(h, t) = \|\mathbf{h}_\perp + \mathbf{r} - \mathbf{t}_\perp\|_2^2.$$


$$\mathbf{h}_\perp = \mathbf{h} - \mathbf{w}_r^\top \mathbf{h} \mathbf{w}_r$$

Solves the problem of 1-to-N,  
N-to-1 and N-to-N relations

# TRANSR

- Authors argue that:

- relations and entities are completely different objects, so they shouldn't be embedded in the same semantic space.
- Even though TransH extends modeling flexibility, it does not perfectly break the restrict of a common semantic space



$$h, t \in \mathbb{R}^k; r \in \mathbb{R}^d$$

$$M_r \in \mathbb{R}^{k \times d}$$

$$\mathbf{h}_r = \mathbf{h}M_r, \quad \mathbf{t}_r = \mathbf{t}M_r.$$

$$f_r(h, t) = \|\mathbf{h}_r + \mathbf{r} - \mathbf{t}_r\|_2^2.$$

	$\langle \text{Head}, \text{Tail} \rangle$
1	$\langle \text{Africa, Congo} \rangle, \langle \text{Asia, Nepal} \rangle, \langle \text{Americas, Aruba} \rangle, \langle \text{Oceania, Federated States of Micronesia} \rangle$
2	$\langle \text{United States of America, Kankakee} \rangle, \langle \text{England, Bury St Edmunds} \rangle, \langle \text{England, Darlington} \rangle, \langle \text{Italy, Perugia} \rangle$
3	$\langle \text{Georgia, Chatham County} \rangle, \langle \text{Idaho, Boise} \rangle, \langle \text{Iowa, Polk County} \rangle, \langle \text{Missouri, Jackson County} \rangle, \langle \text{Nebraska, Cass County} \rangle$
4	$\langle \text{Sweden, Lund University} \rangle, \langle \text{England, King's College at Cambridge} \rangle, \langle \text{Fresno, California State University at Fresno} \rangle, \langle \text{Italy, Milan Conservatory} \rangle$

Basic idea of CTransR: Grouping head-tail pairs into different clusters and learning relation embeddings for each cluster

## CTANSR – CLUSTER-BASE D TRANSR

A UNIQUE VECTOR FOR  
EACH RELATION MIGHT BE  
UNDER-REPRESENTATIVE

1. Obtain entity embeddings  $h$  and  $t$  for all  $(h, t)$  pairs using TransE
2. Compute vector offsets  $(h - t)$  for all training data for each relation  $r$
3. Vector offsets for a certain relation are likely to form multiple clusters
4. Learn a separate relation vector  $r_c$  for each cluster and matrix  $M_r$  for each relation, respectively (Authors seem to assume different clusters within the same relation share a single relation space)

$$\mathbf{h}_{r,c} = \mathbf{hM}_r \quad \mathbf{t}_{r,c} = \mathbf{tM}_r$$

$$f_r(h, t) = \|\mathbf{h}_{r,c} + \mathbf{r}_c - \mathbf{t}_{r,c}\|_2^2 + \alpha \|\mathbf{r}_c - \mathbf{r}\|_2^2,$$

## CTRANSR – CLUSTER-BASED TRANSR

A UNIQUE VECTOR FOR EACH RELATION MIGHT BE UNDER-REPRESENTATIVE

# EXPERIMENT RESULTS

Link Prediction: predicting the missing  $h$  or  $t$  for a relation fact triple  $(h, r, t)$

Data Sets Metric	WN18				FB15K			
	Mean Rank		Hits@10 (%)		Mean Rank		Hits@10 (%)	
	Raw	Filter	Raw	Filter	Raw	Filter	Raw	Filter
Unstructured (Bordes et al. 2012)	315	304	35.3	38.2	1,074	979	4.5	6.3
RESCAL (Nickel, Tresp, and Kriegel 2011)	1,180	1,163	37.2	52.8	828	683	28.4	44.1
SE (Bordes et al. 2011)	1,011	985	68.5	80.5	273	162	28.8	39.8
SME (linear) (Bordes et al. 2012)	545	533	65.1	74.1	274	154	30.7	40.8
SME (bilinear) (Bordes et al. 2012)	526	509	54.7	61.3	284	158	31.3	41.3
LFM (Jenatton et al. 2012)	469	456	71.4	81.6	283	164	26.0	33.1
TransE (Bordes et al. 2013)	263	251	75.4	89.2	243	125	34.9	47.1
TransH (unif) (Wang et al. 2014)	318	303	75.4	86.7	211	84	42.5	58.5
TransH (bern) (Wang et al. 2014)	401	388	73.0	82.3	212	87	45.7	64.4
TransR (unif)	232	219	78.3	91.7	226	78	43.8	65.5
TransR (bern)	238	225	<b>79.8</b>	92.0	<b>198</b>	77	48.2	68.7
CTransR (unif)	243	230	78.9	<b>92.3</b>	233	82	44	66.3
CTransR (bern)	<b>231</b>	<b>218</b>	79.4	<b>92.3</b>	199	<b>75</b>	<b>48.4</b>	<b>70.2</b>


## EXPERIMENT RESULTS

Triple Classification: judging whether a given triple  $(h, r, t)$  is correct

Data Sets	WN11	FB13	FB15K
SE	53.0	75.2	-
SME (bilinear)	70.0	63.7	-
SLM	69.9	85.3	-
LFM	73.8	84.3	-
NTN	70.4	<b>87.1</b>	68.5
TransE (unif)	75.9	70.9	79.6
TransE (bern)	75.9	81.5	79.2
TransH (unif)	77.7	76.5	79.0
TransH (bern)	78.8	83.3	80.2
TransR (unif)	85.5	74.7	81.7
TransR (bern)	<b>85.9</b>	82.5	83.9
CTransR (bern)	85.7	-	<b>84.5</b>

# EXPERIMENT RESULTS

Relation Extraction from Text: Combining results from text-based relation extraction model and knowledge graph embeddings to rank test triples





## MY THOUGHTS

- Training time – Performance Tradeoff
- A single CNN instead of matrix for each relation
- Relation hyperplane vs. relation space
- CTransR is more inspirational

# Gated Graph Sequence Neural Networks

Li, Y., Tarlow, D., Brockschmidt, M., & Zemel, R, ICLR 2016

Presented by Hyoungwook Nam (hn5)

# Abstract

- **Graph-structured data** appears on many domains
- Based on **GNNs** (graph neural network), utilize **GRU** (gated recurrent unit) and extend to output **sequences**
- The result is **flexible**, and better than sequence-based models (e.g. LSTM) if a problem can be **graph-structured**
- State-of-the-art on **bAbI** and **graph algorithm** tasks

# Introduction


Previous approaches:

- Graph feature engineering, Graph neural network (GNN), spectral networks, etc.

Contributions:


- Propose GGS-NN, a **gated GNN** for **sequence output**.
- Show that it is useful for many problems (shortest path, program verification, etc.)

# Graph Neural Network (GNN)



- Propagation model gives node representations (embeddings)
- **Output model**  $g$  provides outputs  $o_v = g(\mathbf{h}_v, l_v)$  per vertex
- Similar to RNN encoder-decoder without attention

# Propagation Model



- $\mathbf{h}_v^{(t)} = f(NBR_v^{(t-1)})$  where  $NBR_v$  is a set of v's **neighbors**
- From initial  $\mathbf{h}_v^{(1)}$ 's, the update **repeats until convergence**

# Gated Graph Neural Network (GG-NN)

- Initialize  $h_v^{(1)}$  with **annotations**  $x_v$  instead of random values
- **GRU-like** propagation model

$$\mathbf{h}_v^{(1)} = [x_v^\top, \mathbf{0}]^\top \quad (1)$$

$$\mathbf{a}_v^{(t)} = \mathbf{A}_{v:}^\top \left[ \mathbf{h}_1^{(t-1)\top} \dots \mathbf{h}_{|\mathcal{V}|}^{(t-1)\top} \right]^\top + \mathbf{b} \quad (2)$$

$$\mathbf{z}_v^t = \sigma \left( \mathbf{W}^z \mathbf{a}_v^{(t)} + \mathbf{U}^z \mathbf{h}_v^{(t-1)} \right) \quad (3)$$

$$\mathbf{r}_v^t = \sigma \left( \mathbf{W}^r \mathbf{a}_v^{(t)} + \mathbf{U}^r \mathbf{h}_v^{(t-1)} \right) \quad (4)$$


$$\widetilde{\mathbf{h}_v^{(t)}} = \tanh \left( \mathbf{W} \mathbf{a}_v^{(t)} + \mathbf{U} \left( \mathbf{r}_v^t \odot \mathbf{h}_v^{(t-1)} \right) \right) \quad (5)$$

$$\mathbf{h}_v^{(t)} = (1 - \mathbf{z}_v^t) \odot \mathbf{h}_v^{(t-1)} + \mathbf{z}_v^t \odot \widetilde{\mathbf{h}_v^{(t)}}. \quad (6)$$

- Output model: Graph-level or node-selection with softmax

$$\mathbf{h}_G = \tanh \left( \sum_{v \in \mathcal{V}} \sigma \left( i(\mathbf{h}_v^{(T)}, \mathbf{x}_v) \right) \odot \tanh \left( j(\mathbf{h}_v^{(T)}, \mathbf{x}_v) \right) \right) \quad o_v = g(\mathbf{h}_v^{(T)}, \mathbf{x}_v)$$


# Adjacency Matrix and Neighborhood



- Adjacency matrix  $A = [A^{(\text{out})}, A^{(\text{in})}]$  for neighborhood updates
- $a_v^{(t)} = A^T [h_1^{(t-1)} \dots h_V^{(t-1)}]^T$  will propagate  $h_{v'}$  of v's neighbors
- $h_v^{(t)} = GRU(a_v^{(t)}, h_v^{(t-1)})$

# Gated Graph Sequence NN (GGS-NN)


- Objective: create an output sequence  $o^{(1)} \dots o^{(k)}$
- RNN-like structure using two GG-NNs  $F_o^{(n)}, F_x^{(n)}$



- Latent (hidden) or observed annotations  $X^{(n)}$ s are possible

# bAbI Task Evaluation Setup

- Symbolic task to graph structured problem



# bAbI + Graph Algorithm Result

- (N): Samples needed for the best result (max 950)

Task	RNN	LSTM	GG-NN
bAbI Task 4	97.3±1.9 (250)	97.4±2.0 (250)	100.0±0.0 (50)
bAbI Task 15	48.6±1.9 (950)	50.3±1.3 (950)	100.0±0.0 (50)
bAbI Task 16	33.0±1.9 (950)	37.5±0.9 (950)	100.0±0.0 (50)
bAbI Task 18	88.9±0.9 (950)	88.9±0.8 (950)	100.0±0.0 (50)

Task	RNN	LSTM	GGS-NNs		
bAbI Task 19	24.7±2.7 (950)	28.2±1.3 (950)	71.1±14.7 (50)	92.5±5.9 (100)	99.0±1.1 (250)
Shortest Path	9.7±1.7 (950)	10.5±1.2 (950)	100.0± 0.0 (50)		
Eulerian Circuit	0.3±0.2 (950)	0.1±0.2 (950)	100.0± 0.0 (50)		



# Program Verification Setup

- Program → Memory Heap → GG-NN → Invariant Logic

```
node* concat(node* a, node* b) {  
    if (a == NULL) return b;  
    node* cur = a;  
    while (cur.next != NULL)  
        cur = cur->next;  
    cur->next = b;  
    return a;  
}
```



GGS-NN


$$\text{ls}(\text{arg1}, \text{NULL}, \lambda t_1 \rightarrow \text{ls}(t_1, \text{NULL}, \top)) * \text{tree}(\text{arg2}, \lambda t_2 \rightarrow \exists e_1. \text{ls}(t_2, e_1, \top) * \text{ls}(e_1, e_1, \top))$$

# Program Verification Result

- Exceeds the previous method with domain-specific feature engineering (89.96% > 89.11%)

Program	Invariant Found
Traverse1	$\text{ls}(\text{lst}, \text{curr}) * \text{ls}(\text{curr}, \text{NULL})$
Traverse2	$\text{curr} \neq \text{NULL} * \text{lst} \neq \text{NULL} * \text{ls}(\text{lst}, \text{curr}) * \text{ls}(\text{curr}, \text{NULL})$
Concat	$\text{a} \neq \text{NULL} * \text{a} \neq \text{b} * \text{b} \neq \text{curr} * \text{curr} \neq \text{NULL}$ $* \text{ls}(\text{curr}, \text{NULL}) * \text{ls}(\text{a}, \text{curr}) * \text{ls}(\text{b}, \text{NULL})$
Copy	$\text{ls}(\text{curr}, \text{NULL}) * \text{ls}(\text{lst}, \text{curr}) * \text{ls}(\text{cp}, \text{NULL})$
Dispose	$\text{ls}(\text{lst}, \text{NULL})$
Insert	$\text{curr} \neq \text{NULL} * \text{curr} \neq \text{elt} * \text{elt} \neq \text{NULL} * \text{elt} \neq \text{lst} * \text{lst} \neq \text{NULL}$ $* \text{ls}(\text{elt}, \text{NULL}) * \text{ls}(\text{lst}, \text{curr}) * \text{ls}(\text{curr}, \text{NULL})$
Remove	$\text{curr} \neq \text{NULL} * \text{lst} \neq \text{NULL} * \text{ls}(\text{lst}, \text{curr}) * \text{ls}(\text{curr}, \text{NULL})$

# Takeaways

- GNNs consist of a **propagation model** to update node representations and an **output model** to compute the outputs
- GG-NN uses a **GRU-like** propagation model and GGS-NN follows the recurrent structure for **sequential outputs**
- They are proven very powerful on tasks like bAbI and program verification which can be **graph-structured**

# Graph CNNs for Semantic Role Labeling

Eddie Huang

marcheggiani-titov-2017-encoding  
"Encoding Sentences with Graph Convolutional Networks for Semantic Role  
Labeling" - Marcheggiani, Diego and Titov, Ivan

20 February, 2020

# Outline

## Main Idea

### Introduction

Semantic Role Labeling (SRL)

Related Work

## Reiterate Main Idea

### Methodology

Syntactic Dependency Graph

Graph Convolutional Neural Networks (GCNs)

Architecture

## Results

## Criticism

# Outline

## Main Idea

### Introduction

Semantic Role Labeling (SRL)

Related Work

### Reiterate Main Idea

## Methodology

Syntactic Dependency Graph

Graph Convolutional Neural Networks (GCNs)

Architecture

## Results

## Criticism

# Main Idea

# Main Idea





Figure 1: Model Architecture

# Main Idea



A new model using graph convolutional neural networks with syntax graphs exceeds previous best models in **semantic role labeling**

Figure 1: Model Architecture

# Outline

Main Idea

## Introduction

Semantic Role Labeling (SRL)

Related Work

Reiterate Main Idea

## Methodology

Syntactic Dependency Graph

Graph Convolutional Neural Networks (GCNs)

Architecture

Results

Criticism

# What is Semantic Role Labeling (SRL)?

# What is Semantic Role Labeling (SRL)?

Want to know "**who did what to whom?**"

# What is Semantic Role Labeling (SRL)?

Want to know "**who did what to whom?**"

## Example

Sequa makes and repairs jet engines

# What is Semantic Role Labeling (SRL)?

Want to know "**who did what to whom?**"

## Example

Sequa makes and repairs jet engines

- ▶ Predicates: makes, repairs
- ▶ Semantic Roles:
  - ▶ Agent: Sequa
  - ▶ Patient: engines

# Why do we want SRL?

## Why do we want SRL?



Figure 2: SRL provides more intermediate features in NLP pipeline

## Related Work

## Related Work

- ▶ Earliest works with RNNs on SRL began in 2008
- ▶ 2014-2017 Modern approaches using LSTMs and Syntactic features
- ▶ A multi-layer Bi-LSTM model made in 2017 was the most state-of-the-art SRL model at the time (created by the same author)

## Basic Components

- ▶ GCNs
- ▶ Syntax Parsing
- ▶ LSTMs
- ▶ Word Embeddings

# Outline

Main Idea

Introduction

Semantic Role Labeling (SRL)

Related Work

Reiterate Main Idea

Methodology

Syntactic Dependency Graph


Graph Convolutional Neural Networks (GCNs)

Architecture

Results

Criticism

## Reiterate Main Idea



A new model using graph convolutional neural networks with syntax graphs exceeds previous best models in **semantic role labeling**

Figure 3: Model Architecture

## Example




Figure 4: An Example (red is what we want to find)

# Outline

Main Idea

Introduction

Semantic Role Labeling (SRL)

Related Work

Reiterate Main Idea

Methodology

Syntactic Dependency Graph

Graph Convolutional Neural Networks (GCNs)

Architecture

Results

Criticism

## Syntactic Dependency Graph

- ▶ Syntax of a language can be represented as a relationship between words rooted at the predicate of a sentence




Figure 5: A syntax dependency graph

## Syntactic Dependency Graph

- ▶ Syntax of a language can be represented as a relationship between words rooted at the predicate of a sentence
- ▶ Edges represent the syntactic relationship between the nodes




Figure 5: A syntax dependency graph

# Role of Syntactic Dependency Graphs

# Role of Syntactic Dependency Graphs




Figure 6: Syntactic dependency occurs between LSTM and GCN

# What are Graph Convolutional Neural Networks (GCNs)?

# What are Graph Convolutional Neural Networks (GCNs)?





Figure 7: Graph Convolutional Neural Network

# What are Graph Convolutional Neural Networks (GCNs)?



GCNs are neural networks that take in a graph (a set of nodes and edges) and **output features for each node**.

Figure 7: Graph Convolutional Neural Network

# How do GCNs compute features for nodes?

## How do GCNs compute features for nodes?

Node features are computed as **non-linear combinations of their neighbors**

## How do GCNs compute features for nodes?

Node features are computed as **non-linear combinations of their neighbors**

$$h_v = \text{ReLU} \left( \sum_{u \in N(v)} (Wx_u + b) \right)$$

- ▶  $x_u$  is a vector representation of node  $u$ .

## How do GCNs compute features for nodes?

Node features are computed as **non-linear combinations of their neighbors**

$$h_v = \text{ReLU} \left( \sum_{u \in N(v)} (Wx_u + b) \right)$$

- ▶  $x_u$  is a vector representation of node  $u$ .

Can stack  $k$  GCN layers to capture dependency between nodes  $k$  hops away ( $k = 1$  was best)

$$h_v^{(k)} = \text{ReLU} \left( \sum_{u \in N(v)} (W^{(k-1)} h_u^{(k-1)} + b^{(k-1)}) \right)$$

# K Layers Captures K-hop dependencies






Figure 8:  $h_v^{(k)} = \text{ReLU} \left( \sum_{u \in N(v)} (W^{(k-1)} h_u^{(k-1)} + b^{(k-1)}) \right)$

# Capturing Edge Information



**Figure 9:** Syntax graphs have directionality and edges have different meanings based on their syntax

# Capturing Edge Information



**Figure 9:** Syntax graphs have directionality and edges have different meanings based on their syntax

**Solution** - Have separate weights for each type of edge

# Capturing Edge Information




Figure 9: Syntax graphs have directionality and edges have different meanings based on their syntax

**Solution** - Have separate weights for each type of edge

$$h_v^{(k)} = \text{ReLU} \left( \sum_{u \in N(v)} (W_{dir(u,v)}^{(k-1)} h_u^{(k-1)} + b_{L(u,v)}^{(k-1)}) \right)$$

- ▶  $dir(u, v) \in \{backward(1), self-loop(2), forward(3)\}$
- ▶  $L(u, v)$  captures both directionality and syntax function

# Weighting Importance to Different Syntax




Figure 10: Some edges are more important than others

# Weighting Importance to Different Syntax




Figure 10: Some edges are more important than others

**Solution** - Use sigmoid to express weighted importance

# Weighting Importance to Different Syntax




Figure 10: Some edges are more important than others

**Solution** - Use sigmoid to express weighted importance

$$g_{u,v}^{(k-1)} = \sigma(h_u^{(k-1)} \cdot \hat{v}_{dir(u,v)}^{(k-1)} + \hat{b}_{L(u,v)}^{(k-1)})$$

# Weighting Importance to Different Syntax




Figure 10: Some edges are more important than others

**Solution** - Use sigmoid to express weighted importance

$$g_{u,v}^{(k-1)} = \sigma(h_u^{(k-1)} \cdot \hat{v}_{dir(u,v)}^{(k-1)} + \hat{b}_{L(u,v)}^{(k-1)})$$

$$h_v^{(k)} = ReLU \left( \sum_{u \in N(v)} g_{u,v}^{(k-1)} \left( W_{dir(u,v)}^{(k-1)} h_u^{(k-1)} + b_{L(u,v)}^{(k-1)} \right) \right)$$

## Final Version of GCN


Node features are computed as a weighted non-linear combination of neighbors within  $k$  hops.

$$h_v^{(k)} = \text{ReLU} \left( \sum_{u \in N(v)} g_{u,v}^{(k-1)} \left( W_{\text{dir}(u,v)}^{(k-1)} h_u^{(k-1)} + b_{L(u,v)}^{(k-1)} \right) \right)$$

### Remark

Similar to a multi-layer perceptron

# Architecture



## Remarks

- ▶ Relies on external syntactic parser and predicate identifier.
- ▶ Layer after the GCN is just another feed-forward network with a softmax for semantic role classification.

Figure 11: Architecture of new model

## LSTMs and GCNs compliment each other




Figure 12: **engines** is physically far away from **makes** but syntactically adjacent to it

LSTMs (RNNs) efficiently capture physically close dependencies.  
GCNs can efficiently capture physically far away dependencies

# Outline

Main Idea

Introduction

Semantic Role Labeling (SRL)

Related Work

Reiterate Main Idea

Methodology

Syntactic Dependency Graph

Graph Convolutional Neural Networks (GCNs)

Architecture

Results

Criticism

# Results

## Precision, recall, and F1 scores for the CoNLL-2009 English and Chinese datasets

System	P	R	F <sub>1</sub>
Lei et al. (2015) (local)	-	-	86.6
FitzGerald et al. (2015) (local)	-	-	86.7
Roth and Lapata (2016) (local)	88.1	85.3	86.7
Marcheggiani et al. (2017) (local)	88.7	86.8	87.7
<b>Ours (local)</b>	<b>89.1</b>	<b>86.8</b>	<b>88.0</b>
Björkelund et al. (2010) (global)	88.6	85.2	86.9
FitzGerald et al. (2015) (global)	-	-	87.3
Foland and Martin (2015) (global)	-	-	86.0
Swayamdipta et al. (2016) (global)	-	-	85.0
Roth and Lapata (2016) (global)	90.0	85.5	87.7
FitzGerald et al. (2015) (ensemble)	-	-	87.7
Roth and Lapata (2016) (ensemble)	90.3	85.7	87.9
<b>Ours (ensemble 3x)</b>	<b>90.5</b>	<b>87.7</b>	<b>89.1</b>

Figure 13: English Results

## Remark

- ▶ Beats previous best results by 0.6% – 1.9%
- ▶  $k = 1$  works best

System	P	R	F <sub>1</sub>
Zhao et al. (2009) (global)	80.4	75.2	77.7
Björkelund et al. (2009) (global)	82.4	75.1	78.6
Roth and Lapata (2016) (global)	83.2	75.9	79.4
<b>Ours (local)</b>	<b>84.6</b>	<b>80.4</b>	<b>82.5</b>

Figure 14: Chinese Results

# Outline

Main Idea

Introduction

Semantic Role Labeling (SRL)

Related Work

Reiterate Main Idea

Methodology

Syntactic Dependency Graph

Graph Convolutional Neural Networks (GCNs)

Architecture

Results

Criticism

## Criticism

Syntactic graph parsing is similar to semantic role labeling because their graph structures look nearly the same. Could probably make at least a decent hand-made algorithm to perform SRL given syntax dependency graph. Would like to see comparison between hand-made algorithm vs. neural net.




Figure 15: SRL and Syntactic are nearly identical